

Story 1	The Brilliant, Untrustworthy Intern
Story Title	The Brilliant, Untrustworthy Intern
Meta Goal	To shift the learner’s mental model of an LLM from a “magic brain” to a powerful but flawed tool that requires constant verification.
Target Audience	Professionals new to generative AI, prone to either excessive trust or dismissive skepticism.
Time	8 minutes (Storytelling: 4 min, CARE walkthrough: 4 min)
Context	AI Literacy, Level 1
Version	
Starting Competence	Unconscious Incompetence (Treats AI output as fact, doesn’t know the risks of hallucination).
Ending Competence	Conscious Incompetence (Understands that AI hallucinates and that all output must be verified, but isn’t yet skilled at it).
Story Objective	Learners walk away with the metaphor “brilliant but untrustworthy intern” as their primary mental model for interacting with LLMs.
Essential Question	How do we responsibly use a tool that is both incredibly powerful and fundamentally unreliable?
Anchor Question	“You wouldn’t let a new intern publish a report without checking their work. Why would you treat this AI any differently?”
Story Arc	Setup: Wow, this thing is a magic brain! It answers anything. Tension: But the answers feel <i>too</i> perfect. Can I trust this? Disruption: The “Untrustworthy Intern” metaphor is introduced, reframing the AI’s confidence as a bug, not a feature. Insight: Oh, my job isn’t to ask questions; it’s to manage this “intern” by giving clear instructions and verifying its work.
Teaching Tools	Slide visuals (“Magic Brain” vs. “Intern”), Live demo of a simple prompt vs. a CARE prompt.
Teaching Situation	Full group lecture, followed by individual exercise on their own laptops.
Tension & Friction	Forcing learners to confront the idea that the confident, well-written text is not a sign of intelligence, but a statistical artifact that can be dangerously wrong.

Story 1	The Brilliant, Untrustworthy Intern
Exercise	“Now that you know how to instruct your ‘intern’ with CARE, let’s try it. Pick a
Hook	complex topic you know, and one you don’t, and see if you can get a good explanation.”
Teaching	Learners experience the dramatic difference in output quality between a simple
Core	prompt and a well-structured CARE prompt. They then feel the unease of needing to verify an answer for a topic they don’t know well, making the competence gap tangible.
